

DB1 Syntax Details

Felix Pojtinger

June 29, 2021

Contents

DB1 Syntax Details	2
Acknowledgements	2
Reset Everything	2
SQL	3
Operators	3
Joins	3
Aliases	4
Limits and Pagination	5
Dates and Intervals	5
Expressions	6
Grouping and Ordering	7
Counting and Sums	7
Inserting	8
Switches	8
Helper Functions	8
Auto-Generated Primary Keys	9
Modifying Columns	9
Virtual Columns	10
Modifying Tables	10
Constraints	10
Types	11
Views	11
Indexes	12
PL/SQL	12
Block Structure	12
Variables	13
Fetching Data	14
Branches and Expressions	14
Switches	15
Labels and Goto	15
Loops	16
Types and Objects	17

Exceptions	17
Cursors	18
Locks	19
Procedures	19
Functions	20
Packages	21
Triggers	22
Maps	24
Arrays	25

DB1 Syntax Details

“so basically i am monkey” - monke, *monkeeee*

Acknowledgements

Most of the following is based on the Oracle Tutorial.

Reset Everything

Run the following to get the commands to drop all tables and their constraints:

```
begin
  for i in (select index_name from user_indexes where index_name not like '%_PK') loop
    execute immediate 'drop index ' || i.index_name;
  end loop;

  for i in (select trigger_name from user_triggers) loop
    execute immediate 'drop trigger ' || i.trigger_name;
  end loop;

  for i in (select view_name from user_views) loop
    execute immediate 'drop view ' || i.view_name;
  end loop;

  for i in (select table_name from user_tables) loop
    execute immediate 'drop table ' || i.table_name || ' cascade constraints';
  end loop;

  execute immediate 'purge recyclebin';
end;
```

Now copy & paste the output into SQL Developer’s SQL worksheet and hit F5.

SQL

Operators

Operator	Description
=	Equality
!=,<>	Inequality
>	Greater than
<	Less than
>=	Greater than or equal to
<=	Less than or equal to
IN	Equal to any value in a list of values
ANY/ SOME/ ALL	Compare a value to a list or subquery. It must be preceded by another operator such as =, >, <.
NOT IN	Not equal to any value in a list of values
[NOT] BETWEEN n and m	Equivalent to [Not] >= n and <= y.
[NOT] EXISTS	Return true if subquery returns at least one row
IS [NOT] NULL	NULL test

Joins

- An **inner join** matches stuff in both tables:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- A **left (outer) join** matches everything in the left tables plus what matches in the right table:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- This **left (outer) join** matches everything that is in the left table and not in the right table:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- A **right (outer) join** matches everything in the right join plus what matches in the left table:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- This **right (outer) join** matches everything that is in the right table and not in the left table:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- A **full (outer) join** merges both tables:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- This **full (outer) join** merges both tables and removes those rows which are in both:

```
select a.id as id_a, a.color as color_a, b.id as id_b, b.color as color_b from palette_
```

- In addition to the on keyword you can also use the using keyword if the PK and FK are the same:

```
select * from orders inner join order_items using(order_id)
```

- You can also use multiple on or using statements:

```
select * from orders inner join order_items using(order_id) inner join customers using
```

- If you use the on keyword, use and for multiples!
- You can also create the Cartesian product:

```
select * from products cross join warehouse;
```

- It is also possible to do a self join:

```
select (w.first_name || ' ' || w.last_name) "Worker", (m.first_name || ' ' || m.last_na
```

- What is the difference between join and union? join merges horizontally (there are more columns than before, maybe also more rows), union merges vertically (there are more rows than before, but the column count stays the same).
- union is similar to T1 | T2 in TypeScript; you can use order by and union to remove duplicates, but note that we have to use select two times:

```
select first_name, last_name, email, 'contact' as role from contacts union select first
```

- union all is similar to T1 & T2 in TypeScript; it keeps duplicates:

```
select last_name from contacts union all select last_name from employees;
```

- Wish to find the difference between two tables? Use intersect:

```
select last_name from contacts intersect select last_name from employees;
```

- Wish to subtract one table from another table? Use minus:

```
select last_name from contacts minus select last_name from employees;
```

Aliases

- You can alias long column names with select mylongname as name from contacts or just select mylongname name from contacts. The as keyword is optional. Full-text column names are supported by enclosing in ““. as can also format strings: select first_name || ' ' || last_name as "Name" from employees; yields Alice, Bob and System.

- You can also create a table alias (using `from employees e`), but you CAN'T USE the `as` keyword.

Limits and Pagination

- The Oracle equivalent of filter is `fetch n next rows only`: `select * from products order by list_price desc fetch next 5 rows only;`
- You may also use the `fetch next n percent rows only`:
`select * from inventories order by quantity desc fetch next 10 percent rows only;`
- Filtering by for example a quantity, and you only want the first 10 “condition matches”? Use `fetch n next rows with ties`:
`select * from inventories order by quantity desc fetch next 5 rows with ties;`
- Need Pagination? Use `offset`:
`select * from products order by standard_cost desc offset 10 rows fetch next 10 rows on`

Dates and Intervals

- Want to extract a year from a date? Use `extract`:
`select * from orders where status = 'Shipped' and extract(year from order_date) = 2017`
- Want to get the current date? Use `current_date`:
`select current_date from dual;`
- The `to_char` function can convert dates (and timestamps) to chars:
`select to_char(sysdate, 'YYYY-MM-DD') from dual;`
- The `to_date` function can convert chars to dates:
`select to_date('2021-01-12', 'YYYY-MM-DD') from dual;`
- Alternatively, the date literal uses the YYYY-MM-DD format and does not require format specs:
`select date '1969-04-20' from dual;`
- You can get the current date with `sysdate`:
`select localtimestamp from dual;`
- You can get the current date & time with `datelocaltimestamp`:
`select localtimestamp from dual;`
- The current time zone is available with `sessiontimezone`:
`select sessiontimezone from dual (yields Europe/Berlin);`

- The `timestamp` literal uses the YYYY-MM-DD HH24:MI:SS.FF format:

```
select timestamp '1969-04-20 00:00:00.00' from dual;
```

- You may also append the timezone (But keep in mind that timestamp with time zone is the column type in this case):

```
select timestamp '1969-04-20 00:00:00.00 Europe/Berlin' from dual;
```

- The `interval` literal can be used to create intervals:

```
select interval '9' day from dual, select interval '9' month from dual, select interval
```

- Using the `months_between` function, the count of months between two dates can be computed.

Expressions

- Only single quotes are supported.
- Comparisons are done with `=`, `NOT =`.
- It also supports full expression evaluation:

```
select product_name as "Product Name", list_price - standard_cost as "Gross Profit" from pr
```

- You can use `()` in `where` clauses to prioritize:

```
select * from orders where (
status = 'Canceled' or status = 'Pending' ) and customer_id = 44
order by order_date;
```

- The `in` keyword is a useful tool for sub collections and subqueries:

```
- select * from orders where salesman_id in (54, 55, 56)
  order by order_id;
- select * from orders where salesman_id not in (54, 55,
  56) order by order_id; (you can use not)
- select * from employees where employee_id in ( select
  distinct salesman_id from orders where status = 'Canceled'
  ) order by first_name; (you can of course also use not)
```

- `between` can also be used for dates:

```
select * from orders where order_date between date '2016-12-01' and date '2016-12-31'
```

- ... `like '%Asus%'` (note the `'s`) is basically a full-text search.

- Some examples of `like` (you can use not for all of them):

```
- select * from contacts where last_name like 'St%'
- select * from contacts where last_name like '%St'
- select * from contacts where last_name like '%St%'
- select * from contacts where last_name like 'Po_tinger'
  (_ matches any one character)
- select * from contacts where lower(last_name) like 'st%'
```

- `select * from contacts where upper(last_name) like 'st%'`
- `select * from discounts where discount_message like '%%'` (returns everything)
- `select * from discounts where discount_message like '%%' escape '!' (returns everything that includes the string '%')`
- You can compare against null with `is null` (`= NULL` does not work). You can negate with `not`.

Grouping and Ordering

- You can use functions like `upper` and dates when ordering.
- The `group by` keyword can be used to find unique data:
`select status from orders group by status;`
- By combining `group by` with `count` you can count the amount of unique data:
`select status, count (*) from orders group by status;`
- `group by` can also be used with the `where` keyword:
`select name, count(*) as "Shipped Orders" from orders inner join customers using(customer_id) where status = 'Shipped' group by name;`
- `where` can NOT APPEAR AFTER `group by`; use the `having` keyword instead.
- The `having` keyword enables you to filter like with `where`, but after the `group by` keyword like so:
`select status from orders where extract(year from order_date) > '2016' group by status;`
- Multiple order by statements? First ordered by first statement, then “sub-ordered” by the second (last name the same -> now first name is evaluated).
- Want to have nulls first when ordering? Use `nulls first` or `nulls last` as the suffix.
- Removal of duplicates is done with `select distinct`. When multiple columns are being selected, use only one `distinct` keyword at the start. Multiple nulls are filtered (Null = Null).

Counting and Sums

- You can count the amount of rows with the `count()` function:
`select count(*) from products`
- The `sum` function can be used to calculate a total:
`select sum(unit_price * quantity) from order_items;`

- It can also be used to calculate a total per row (the group by order_id part is required; group by order_value does not work):

```
select order_id, sum(unit_price * quantity) as order_value from order_items group by order_id
```

Inserting

- It is a good idea to always specify the columns when inserting:

```
insert into discounts(discount_name, amount, start_date, expired_date) values ('Summer', 10, '2017-01-01', '2017-06-30')
```

- You can also “insert from select” using insert into:

```
insert into sales(customer_id, product_id, order_date, total) select customer_id, product_id, order_date, total from order_items
```

- It's even possible to “create a table from select” using create table x as, basically coping its schema (where 1 = 0 skips copying the rows):

```
create table sales_2017 as select * from sales where 1 = 0;
```

- Using insert all, it is possible to insert multiple rows at once (note the lack of commas between the into keywords. Here, the subquery is ignored/a placeholder.):

```
insert all into fruits (fruit_name, color) values ('Apple', 'Red') into fruits (fruit_name, color) select fruit_name, color from fruits
```

- You can also use conditions based on the subquery (insert first is the equivalent of a switch case.):

```
insert all when amount < 10000 then into small_orders when amount >= 10000 then into big_orders select * from orders
```

Switches

- Using case it is possible to create if/else constructs:

```
select product_name, list_price, case category_id when 1 then round(list_price * 0.05, 2) else list_price end as discounted_price from products
```

- case is also useful for conditional grouping:

```
select * from locations order by country_id, case country_id when 'US' then state else country end
```

- case also evaluates to an expression, so you can use it for conditional updates:

```
update products set list_price = case when list_price < 20 then 30 else 50 end where list_price < 50
```

Helper Functions

- You can extract substrings with substr: select substr('Alex', 1, 1) from dual;
- Stuff like select upper('uwu') from dual can come in handy.
- Using round it is possible to round numbers (returns 5.23):


```
select round(5.234234234234, 2) from dual;
```

- You can use `replace` to replace strings:

```
update accounts set phone = replace(phone, '+1-', '');
```

- You can use the `floor`, `round` and `ceil` functions to get rounded values.

Auto-Generated Primary Keys

- generated by default as identity is quite useful for auto-incrementing columns such as PKs:

```
create table persons ( person_id number generated by default as identity, first_name va
```

- generated always as identity is the same but does not allow setting it manually.

Modifying Columns

- You can use `desc mytable` to show the schema for a table.

- `alter table` can be used to add columns using `add`:

```
alter table persons add birthdate date not null;
```

- You can also add multiples at once (note that there is no column keyword):

```
alter table persons add ( phone varchar2(20), email varchar2(100) )
```

- `modify` can change the column type (note that there is no column keyword):

```
alter table persons modify birthdate date null;
```

- `drop column` can be used to remove a column

```
alter table persons drop column birthdate;
```

- `rename column` can be used to rename a column:

```
alter table persons rename column first_name to forename;
```

- `rename to` can be used to rename a table:

```
alter table persons rename to people;
```

- `rename promotions to promotions_two` is an alternative syntax.

- You can use the `default` keyword to set a default value:

```
alter table accounts add status number(1,0) default 1 not null.
```

- A more efficient logical version of `drop column` is `set unused column`:

```
alter table suppliers set unused column fax;
```

- You can now drop it using:

```
alter table suppliers drop unused columns;
```

- If you want to physically drop a column, use drop:

```
alter table suppliers drop (email, phone);
```

Virtual Columns

- You can create virtual columns in regular tables without using views with `alter table x add ... as` (note the required `(` after the `as` keyword):

```
alter table parts add (capacity_description as ( case when capacity <= 8 then 'Small' w
```

- The size of a `varchar2` is adjustable afterwards (note that this checks if any current `varchar2`s are larger than the new size and fails if they are.):

```
alter table persons modify first_name varchar2(255);
```

Modifying Tables

- You can drop a table with `drop table`:

```
drop table people;
```

- Appending `purge` clears the recycle bin; appending `cascade constraints` drop all related constraints.

- You can clear a table using `truncate table`:

```
truncate table customers_copy;
```

- The same limitations as with `drop table` concerning constraints apply, so appending `cascade (WITHOUT constraints)` drops all related ones.

- You can clear the recycle bin with:

```
purge recyclebin;
```

Constraints

- It is possible to add constraints (any constraints, a primary key in this example) after creating a table with `add constraint`:

```
alter table purchase_orders add constraint purchase_orders_order_id_pk primary key(orde
```

- You may remove a constraint with `drop constraint`:

```
alter table purchase_orders drop constraint purchase_orders_order_id_pk;
```

- Instead of removing it, you can also use `disable constraint`:

```
alter table purchase_orders disable constraint purchase_orders_order_id_pk;
```

- And re-enable it with `enable constraint`:

```
alter table purchase_orders enable constraint purchase_orders_order_id_pk;
```

- You can also add foreign key constraints:

```
alter table suppliers add constraint suppliers_supplier_groups_fk foreign key(group_id)
```

- Using a check constraint, arbitrary expressions can be evaluated:

```
alter table parts add constraint check_buy_price_positive check(buy_price > 0);
```

- A unique constraint prevents unwanted duplicates:

```
alter table clients add constraint unique_clients_phone unique(phone);
```

- With a not null constraint, fuzzy logic can be avoided; it is however best to define nullable fields at schema creation, as the syntax differs from the add constraint/drop constraint logic above:

```
alter table clients modify ( 7 phone not null );
```

- You can remove them by modifying it to null explicitly:

```
alter table clients modify ( phone null );
```

Types

- You can create a number within a range: `number(1,0)`.
- The `number` type is used for all types of numbers by specifying precision and scale: `number(6)` (or `number(6,0)`) is a signed integer fitting 6 digits, `number(6,2)` is a float with two digits precision. The DB doesn't just cut off numbers, it rounds them.
- The float type can be emulated by the number type, i.e. `float(2)` is equal to `number(38,2)`. The argument is in bits instead of digits though.
- The `lengthdb` function can be used to get the length of field in bytes.
- The char type has a fixed length: name `char(10)` or name `char(10 bytes)`, meaning that a char always takes up the amount of bytes set. `nchar` is the same but UTF-8 or UTF-16 any doesn't take bytes.
- The `varchar2` type also takes an argument for the length in bytes, which in ASCII corresponds to the amount of characters. `nvarchar2` is the same but UTF-8 or UTF-16 and doesn't take bytes.

Views

- You can create a view with `create view x as select ...`:

```
create view employees_years_of_service as select employee_id, first_name || ' ' || last
```

- If used with `create or replace view`, upserts are possible.
- By appending with `read only`, you can prevent data modifications:

```
create or replace view employees_years_of_service as select employee_id, first_name ||
```

- `drop view x` removes the view.

- Deletions and updates on views are usually fine, but inserts can often be not that useful due to fields being excluded from the view; see `instead of triggers` later on for a solution;
- Subqueries can be used in selects:


```
select * from ( select * from products) where list_price < 100;
```
- They can also be used in updates:


```
update ( select list_price from products ) set list_price = list_price * 1.5;
```

Indexes

- You can create an index with `create index`:


```
create index members_last_name on members(last_name);
```
- You can also create an index spanning multiple columns:


```
create index members_full_name on members(first_name, last_name);
```
- You can drop an index with `drop index`:


```
drop index members_full_name;
```

PL/SQL

Block Structure

- Block structure:


```
declare
  -- declarations
begin
  -- your logic
exception
  -- exception handling
end;
```
- The most simple example is as follows:


```
begin
  dbms_output.put_line('Hello World!');
end;
```
- Use `put_line` from the `dbms_output` package to print to stdout.
- You can use the `declare` section for variables:


```
declare
  message varchar(255) := 'Hello, World!';
begin
```

```

        dbms_output.put_line(message);
    end;

```

- The `exception` block is used to handle exceptions, for example `zero_divide` for divisions by zero (`when others then` handles unexpected other exceptions):

```

declare
    result number;
begin
    result := 1/0;

    exception
        when zero_divide then
            dbms_output.put_line(sqlerrm);
        when others then
            dbms_output.put_line('An unexpected error occurred: ' || sqlerrm);
end;

```

- You always have to specify an execution section; use `null` for a no-op:

```

declare
begin
    null;
end;

```

- You can use `--` for single line comments and `/*` for multi line comments.

Variables

- PL/SQL extends SQL by adding a boolean type (which can have the values `true`, `false` and `null`).
- Variables need not be given a value at declaration if they are nullable:

```

declare
    total_sales number(15,2);
    credit_limit number(10,0);
    contact_name varchar2(255);
begin
    null;
end;

```

- You can use `default` as an alternative to the `:=` operator when assigning variables in the declaration section. DO NOT use `=` when assignment, even re-assignment also uses `:=`.
- If a variable is defined as not null, it can't take a string of length 0:

```

declare
    shipping_status varchar2(25) not null := 'shipped';

```

```

begin
    shipping_status := ''; -- You need to specify any string != ''
end;

```

- Constants are created with the `constant` keyword and forbid reassignment:

```

declare
    price constant number := 10;
begin
    price := 20; -- Will throw an exception
end;

```

Fetching Data

- Use `select ... into` to fetch data into variables; %TYPE infers the type of a column:

```

declare
    customer_name customers.name%TYPE;
    customer_credit_limit customers.credit_limit%TYPE;
begin
    select
        name, credit_limit
    into
        customer_name, customer_credit_limit
    from customers where customer_id = 38;

    dbms_output.put_line(customer_name || ': ' || customer_credit_limit);
end;

```

Branches and Expressions

- `if ... then ... end if` can be used for branching:

```

declare
    sales number := 20000;
begin
    if sales > 10000 then
        dbms_output.put_line('Lots of sales!');
    end if;
end;

```

- Inline expressions are also supported:

```

large_sales := sales > 10000

```

- Booleans need not be compared with `my_bool = true`, a simple `if my_bool then` is fine.

- `elseif ... then` is NOT valid syntax; `elsif ... then` is valid syntax.
- Statements may also be nested:

```
declare
    sales number := 20000;
begin
    if sales > 10000 then
        if sales > 15000 then
            dbms_output.put_line('A new sales record!');
        else
            dbms_output.put_line('Lots of sales!');
        end if;
    end if;
end;
```

Switches

- You may use the `case` keyword for switch cases:

```
declare
    grade char(1);
    message varchar2(255);
begin
    grade := 'A';

    case grade
        when 'A' then
            message := 'Excellent';
        when 'B' then
            message := 'Great';
        when 'C' then
            message := 'Good';
        when 'D' then
            message := 'Fair';
        when 'F' then
            message := 'Poor';
        else
            raise case_not_found;
    end case;

    dbms_output.put_line(message);
end;
```

Labels and Goto

- A label/goto equivalent is also available:

```

begin
    goto do_work;
    goto goodbye;

    <<do_work>>
    dbms_output.put_line('mawahaha');

    <<goodbye>>
    dbms_output.put_line('Goodbye!');
end;

```

Loops

- The equivalent of the `while` loop is the `loop`. `exit/continue` prevents an infinite loop:

```

declare
    i number := 0;
begin
    loop
        i := i + 1;

        dbms_output.put_line('Iterator: ' || i);

        if i >= 10 then
            exit;
        end if;
    end loop;

    dbms_output.put_line('Done!');
end;

```

- For loops can be done using the `for i in 0..100 loop ... end loop` syntax:

```

begin
    for i in 0..100 loop
        dbms_output.put_line(i);
    end loop;
end;

```

- While loops work as you'd expect; but also require the `loop` keyword:

```

declare
    i number := 0;
begin
    while i <= 100 loop
        dbms_output.put_line(i);
    end loop;
end;

```



```

        i := i + 1;
    end loop;
end;

```

Types and Objects

- You can also use %ROWTYPE to infer the type of a row and select an entire row at once:

```

declare
    customer customers%ROWTYPE;
begin
    select * into customer from customers where customer_id = 100;

    dbms_output.put_line(customer.name || '/' || customer.website);
end;

```

- It is also possible to use OOP-style object/row creation thanks to %ROWTYPE:

```

declare
    person persons%ROWTYPE;

begin
    person.person_id := 1;
    person.first_name := 'John';
    person.last_name := 'Doe';

    insert into persons values person;
end;

```

Exceptions

- You can create custom exceptions:

```

declare
    e_credit_too_high exception;
    pragma exception_init(e_credit_too_high, -20001);
begin
    if 10000 > 1000 then
        raise e_credit_too_high;
    end if;
end;

```

- If you want to raise a custom exception, use raise_application_error:

```

declare
    e_credit_too_high exception;

```

```

        pragma exception_init(e_credit_too_high, -20001);
begin
    raise_application_error(-20001, 'Credit is too high!');
end;

```

- Using `sqlcode` and `sqlerrm` you can get the last exception's code/error message.

Cursors

- Using cursors, you can procedurally process data:

```

declare
    cursor sales_cursor is select * from sales;
    sales_record sales_cursor%ROWTYPE;
begin
    update customers set credit_limit = 0;

    open sales_cursor;

    loop
        fetch sales_cursor into sales_record;
        exit when sales_cursor%NOTFOUND;

        update
            customers
        set
            credit_limit = extract(year from sysdate)
        where
            customer_id = sales_record.customer_id;
    end loop;

    close sales_cursor;
end;

```

- Complex exit logic can be avoided using the `for ... loop`:

```

declare
    cursor product_cursor is select * from products;
begin
    for product_record in product_cursor loop
        dbms_output.put_line(product_record.product_name || ': $' || product_record.list_price);
    end loop;
end;

```

- Cursors can also have parameters:

```

declare
    product_record products%rowtype;

```

```

    cursor
    product_cursor (
        low_price number := 0,
        high_price number := 100
    )
    is
        select * from products where list_price between low_price and high_price;
begin
    open product_cursor(50, 100);

    loop
        fetch product_cursor into product_record;
        exit when product_cursor%notfound;

        dbms_output.put_line(product_record.product_name || ': $' || product_record.list_price);
    end loop;

    close product_cursor;
end;

```

Locks

- The DB can also lock fields for safe multiple access:

```

declare
    cursor customers_cursor is select * from customers for update of credit_limit;
begin
    for customer_record in customers_cursor loop
        update customers set credit_limit = 0 where customer_id = customer_record.customer_id;
    end loop;
end;

```

Procedures

- You can create procedures, which are comparable to functions:

```

create or replace procedure
print_contact(customer_id_arg number)
is
    contact_record contacts%rowtype;
begin
    select * into contact_record from contacts where customer_id = customer_id_arg;

    dbms_output.put_line(contact_record.first_name || ' ' || contact_record.last_name);
end;

```

- These procedures can then be executed:

```
begin
    print_contact(50);
end;
```

- Or, without PL/SQL:

```
exec print_contact(50);
```

- Once a procedure is no longer needed, it can be removed with drop procedure:

```
drop procedure print_contact;
```

- It is also possible to infer a row type using sys_refcursor and return rows with dbms_sql.return_result:

```
create or replace procedure
    get_customer_by_credit(min_credit number)
as
    customer_cursor sys_refcursor;
begin
    open customer_cursor for select * from customers where credit_limit > min_credit;

    dbms_sql.return_result(customer_cursor);
end;
```

- You can now call it:

```
exec get_customer_by_credit(50);
```

Functions

- Functions are similar, but require returning a value:

```
create or replace function
    get_total_sales_for_year(year_arg integer)
return number
is
    total_sales number := 0;
begin
    select sum(unit_price * quantity) into total_sales
    from order_items
    inner join orders using (order_id)
    where status = 'Shipped'
    group by extract(year from order_date)
    having extract(year from order_date) = year_arg;

    return total_sales;
end;
```

- You can call them from PL/SQL:

```

declare
    total_sales number := 0;
begin
    total_sales := get_total_sales_for_year(2017);

    dbms_output.put_line('Sales for 2017: ' || total_sales);
end;

```

- And remove them with drop function:

```

drop function get_total_sales_for_year;

```

Packages

- Packages can be used to group function “interfaces” and variables:

```

create or replace package order_management
as
    shipped_status constant varchar(10) := 'Shipped';
    pending_status constant varchar(10) := 'Pending';
    cancelled_status constant varchar(10) := 'Canceled';

    function get_total_transactions return number;
end order_management;

```

- You can now access the variables in the package with .:

```

begin
    dbms_output.put_line(order_management.shipped_status);
end;

```

- In order to use functions in a package, you then have to create a package body, implementing it:

```

create or replace package body order_management
as
    function get_total_transactions return number
    is
        total_transactions number;
    begin
        select sum(unit_price) into total_transactions from orders inner join order_items
        on orders.order_id = order_items.order_id;

        return total_transactions;
    end;
end;

```

- You can now access the functions in the package with .:

```

select
    order_management.get_total_transactions() as total_transactions

```

```

from
    dual;

```

- And the same is possible from PL/SQL:

```

begin
    dbms_output.put_line(order_management.get_total_transactions());
end;

```

- You can drop a package with `drop package` and a package body with `drop package body`:

```

drop package body order_management;
drop package order_management;

```

Triggers

- Triggers follow a similar structure as procedures:

```

declare
    -- declarations
begin
    -- your logic
exception
    -- exception handling
end;

```

- Using triggers, you can for example create a manual log after operations with `after update` or `delete` on ...:

```

create or replace trigger customers_audit_trigger
    after update or delete
    on customers
    for each row
declare
    transaction_type varchar2(10);
begin
    transaction_type := case
        when updating then 'update'
        when deleting then 'delete'
    end;

    insert into audits(
        table_name,
        transaction_name,
        by_user,
        transaction_date
    ) values (
        'customers',

```

```

        transaction_type,
        user,
        sysdate
    );
end;

```

- Thanks to before update of ... on ..., it is also possible to do more complex checks before inserting:

```

create or replace trigger customers_credit_trigger
    before update of credit_limit
    on customers
declare
    current_day number;
begin
    current_day := extract(day from sysdate);

    if current_day between 28 and 31 then
        raise_application_error(-20100, 'Locked at the end of the month');
    end if;
end;

```

- In combination with when, new (not available in delete statements) and old (not available in insert statements), it is also possible to check based on the previous & current values:

```

create or replace trigger customers_credit_limit_trigger
    before update of credit_limit
    on customers
    for each row
    when (new.credit_limit > 0)
begin
    if :new.credit_limit >= 2*old.credit_limit then
        raise_application_error(-20101, 'The new credit cannot be more than double the
    end if;
end;

```

- Using instead of triggers and returning ... into ..., you can also use views to safely insert into multiple tables:

```

create or replace trigger create_customer_trigger
    instead of insert on customers_and_contacts
    for each row
declare
    current_customer_id number;
begin
    insert into customers(
        name,
        address,

```

```

        website,
        credit_limit
    ) values (
        :new.name,
        :new.address,
        :new.website,
        :new.credit_limit
    ) returning customer_id into current_customer_id;

insert into contacts(
    first_name,
    last_name,
    email,
    phone,
    customer_id
) values (
    :new.first_name,
    :new.last_name,
    :new.email,
    :new.phone,
    current_customer_id
);
end;
```

- You can enable/disable a trigger with `alter trigger ... disable/enable`:

```
alter trigger create_customer_trigger disable;
```
- And completely remove it with `drop trigger`:

```
drop trigger create_customer_trigger;
```
- It is also possible to enable/disable all triggers of a table with `alter table ... enable/disable all triggers`:

```
alter table customers enable all triggers;
```

Maps

- Maps are also possible in PL/SQL using `table of`:

```

declare
    type country_capitals_type
    is table of varchar2(100)
    index by varchar2(50);

    country_capitals country_capitals_type;
begin
    country_capitals('China') := 'Beijing';
```



```

country_capitals('EU') := 'Brussels';
country_capitals('USA') := 'Washington';
end;

```

- You can use `mymap.first` and `mymap.next` to iterate:

```

declare
    type country_capitals_type
        is table of varchar2(100)
        index by varchar2(50);

country_capitals country_capitals_type;
current_country varchar2(50);
begin
    country_capitals('China') := 'Beijing';
    country_capitals('EU') := 'Brussels';
    country_capitals('USA') := 'Washington';

    current_country := country_capitals.first;

    while current_country is not null loop
        dbms_output.put_line(current_country || ': ' || country_capitals(current_country));

        current_country := country_capitals.next(current_country);
    end loop;
end;

```

Arrays

- Using `varray`, it is also possible to create arrays:

```

declare
    type names_type is varray(255) of varchar2(20) not null;

names names_type := names_type('Albert', 'Jonathan', 'Judy');
begin
    dbms_output.put_line('Length before append: ' || names.count);

    names.extend;

    names(names.last) := 'Alice';

    dbms_output.put_line('Length after append: ' || names.count);

    names.trim;

    dbms_output.put_line('Length after trim: ' || names.count);

```

```
names.trim(2);

dbms_output.put_line('Length after second trim: ' || names.count);

names.delete;

dbms_output.put_line('Length after delete: ' || names.count);
end;
```